

131HW5

Scott Shang (8458655)

May 15, 2022

Question1

```
library("tidyverse")
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.6    v purrr    0.3.4
## v tibble  3.1.7    v dplyr    1.0.9
## v tidyr   1.2.0    v stringr  1.4.0
## v readr   2.1.2    v forcats  0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library("tidymodels")
```

```
## -- Attaching packages ----- tidymodels 0.2.0 --
```

```
## v broom      0.8.0    v rsample      0.1.1
## v dials      0.1.1    v tune         0.2.0
## v infer      1.0.0    v workflows    0.2.6
## v modeldata  0.1.1    v workflowsets 0.2.1
## v parsnip    0.2.1    v yardstick    0.0.9
## v recipes    0.2.0
```

```
## -- Conflicts ----- tidymodels_conflicts() --
```

```
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed()  masks stringr::fixed()
## x dplyr::lag()       masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step()   masks stats::step()
## * Use suppressPackageStartupMessages() to eliminate package startup messages
```

```
library("dplyr")
library("yardstick")
library(readr)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
```

```
library(discrim)
```

```
##
## Attaching package: 'discrim'

## The following object is masked from 'package:dials':
##
##     smoothness
```

```
library(poissonreg)
library(corr)
library(klaR)
```

```
## Loading required package: MASS
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##     select
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(knitr)
library(MASS)
library(ggplot2)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack

## Loaded glmnet 4.1-4
```

```
library(janitor)
```

```
##
```

```
## Attaching package: 'janitor'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      chisq.test, fisher.test
```

```
pkm=read_csv('Pokemon.csv')
```

```
## Rows: 800 Columns: 13
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (3): Name, Type 1, Type 2
```

```
## dbl (9): #, Total, HP, Attack, Defense, Sp. Atk, Sp. Def, Speed, Generation
```

```
## lgl (1): Legendary
```

```
##
```

```
## i Use 'spec()' to retrieve the full column specification for this data.
```

```
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
pkm
```

```
## # A tibble: 800 x 13
```

```
##   # Name Type 1 Type 2 Total HP Attack Defense Sp. Atk Sp. Def
##   <dbl> <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1 Bulba~ Grass Poison 318 45 49 49 65 65
## 2 2 Ivysa~ Grass Poison 405 60 62 63 80 80
## 3 3 Venus~ Grass Poison 525 80 82 83 100 100
## 4 3 Venus~ Grass Poison 625 80 100 123 122 120
## 5 4 Charm~ Fire <NA> 309 39 52 43 60 50
## 6 5 Charm~ Fire <NA> 405 58 64 58 80 65
## 7 6 Chari~ Fire Flying 534 78 84 78 109 85
## 8 6 Chari~ Fire Dragon 634 78 130 111 130 85
## 9 6 Chari~ Fire Flying 634 78 104 78 159 115
## 10 7 Squir~ Water <NA> 314 44 48 65 50 64
## # ... with 790 more rows, and 3 more variables: Speed <dbl>, Generation <dbl>,
## # Legendary <lgl>
```

```
pkm=clean_names(pkm)
```

```
pkm
```

```
## # A tibble: 800 x 13
```


```
##   number name type_1 type_2 total hp attack defense sp_atk sp_def speed
##   <dbl> <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1 Bulbasaur Grass Poison 318 45 49 49 65 65 45
## 2 2 Ivysaur Grass Poison 405 60 62 63 80 80 60
## 3 3 Venusaur Grass Poison 525 80 82 83 100 100 80
## 4 3 Venusaur~ Grass Poison 625 80 100 123 122 120 80
```

```
## 5      4 Charmand~ Fire  <NA>    309    39    52    43    60    50    65
## 6      5 Charmele~ Fire  <NA>    405    58    64    58    80    65    80
## 7      6 Charizard Fire  Flying    534    78    84    78   109    85   100
## 8      6 Charizar~ Fire  Dragon    634    78   130   111   130    85   100
## 9      6 Charizar~ Fire  Flying    634    78   104    78   159   115   100
## 10     7 Squirtle  Water  <NA>    314    44    48    65    50    64    43
## # ... with 790 more rows, and 2 more variables: generation <dbl>,
## #   legendary <lgl>
```

After using `clean_names()` on our data, it returns name with only lowercase letter, with `_` as a separator, and convert symbol “#” to “number”. It is helpful because it cleans up the names of variables.

Question2

```
ggplot(pkm, aes(x = type_1)) +
  geom_bar()
```



There are 18 classes of the outcome. The flying type has the least numbers of Pokémon. Dark, dragon, fairy, fighting, ghost, ground, ice, poison and steel types also have less pokemon comparing to others.

```
pkm=filter(pkm,type_1 %in% c("Bug","Fire","Grass","Normal","Water","Psychic"))
pkm
```

```
## # A tibble: 458 x 13
```

```
##   number name      type_1 type_2 total    hp attack defense sp_atk sp_def speed
```

```
##      <dbl> <chr>      <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      1 Bulbasaur Grass Poison 318 45 49 49 65 65 45
## 2      2 Ivysaur  Grass Poison 405 60 62 63 80 80 60
## 3      3 Venusaur Grass Poison 525 80 82 83 100 100 80
## 4      3 Venusaur~ Grass Poison 625 80 100 123 122 120 80
## 5      4 Charmand~ Fire <NA> 309 39 52 43 60 50 65
## 6      5 Charmele~ Fire <NA> 405 58 64 58 80 65 80
## 7      6 Charizard Fire Flying 534 78 84 78 109 85 100
## 8      6 Charizar~ Fire Dragon 634 78 130 111 130 85 100
## 9      6 Charizar~ Fire Flying 634 78 104 78 159 115 100
## 10     7 Squirtle Water <NA> 314 44 48 65 50 64 43
## # ... with 448 more rows, and 2 more variables: generation <dbl>,
## #   legendary <lgl>
```

```
pkm$type_1=as.factor(pkm$type_1)
pkm$legendary=as.factor(pkm$legendary)
pkm$generation=as.factor(pkm$generation)
pkm
```

```
## # A tibble: 458 x 13
##   number name      type_1 type_2 total   hp attack defense sp_atk sp_def speed
##   <dbl> <chr>      <fct> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      1 Bulbasaur Grass Poison 318 45 49 49 65 65 45
## 2      2 Ivysaur  Grass Poison 405 60 62 63 80 80 60
## 3      3 Venusaur Grass Poison 525 80 82 83 100 100 80
## 4      3 Venusaur~ Grass Poison 625 80 100 123 122 120 80
## 5      4 Charmand~ Fire <NA> 309 39 52 43 60 50 65
## 6      5 Charmele~ Fire <NA> 405 58 64 58 80 65 80
## 7      6 Charizard Fire Flying 534 78 84 78 109 85 100
## 8      6 Charizar~ Fire Dragon 634 78 130 111 130 85 100
## 9      6 Charizar~ Fire Flying 634 78 104 78 159 115 100
## 10     7 Squirtle Water <NA> 314 44 48 65 50 64 43
## # ... with 448 more rows, and 2 more variables: generation <fct>,
## #   legendary <fct>
```

Question3

```
set.seed(1234)
pkm_split=initial_split(pkm,prop=0.70,strata=type_1)
train=training(pkm_split)
test=testing(pkm_split)
dim(pkm)*0.7
```

```
## [1] 320.6 9.1
```

```
dim(train)
```

```
## [1] 318 13
```

Yes, the training and test sets have the desired number of observations.

```
folds=vfold_cv(train,v=5,strata=type_1)
folds
```

```
## # 5-fold cross-validation using stratification
## # A tibble: 5 x 2
##   splits      id
##   <list>      <chr>
## 1 <split [252/66]> Fold1
## 2 <split [253/65]> Fold2
## 3 <split [253/65]> Fold3
## 4 <split [256/62]> Fold4
## 5 <split [258/60]> Fold5
```

Stratifying the folds is useful because it makes sure that the folds are representative of the whole data set.

Question4

```
rcp=recipe(type_1~legendary+generation+sp_atk+attack+speed+defense+hp+sp_def,data=train) %>%
  step_dummy(legendary) %>%
  step_dummy(generation) %>%
  step_normalize(all_predictors())
```

Question5

```
reg=multinom_reg(mixture=tune(),penalty=tune()) %>%
  set_mode("classification") %>%
  set_engine("glmnet")

wf=workflow() %>%
  add_recipe(rcp) %>%
  add_model(reg)

grid<-grid_regular(penalty(range = c(-5, 5)),mixture(),levels = 10)
grid
```


```
## # A tibble: 100 x 2
##       penalty mixture
##       <dbl>   <dbl>
## 1      0.00001      0
## 2      0.000129     0
## 3      0.00167      0
## 4      0.0215       0
## 5      0.278        0
## 6      3.59         0
## 7     46.4          0
## 8     599.          0
## 9     7743.         0
## 10 100000           0
## # ... with 90 more rows
```

We will be fitting 500 models, since we have 10 levels of penalty, 10 levels of mixture, and 5 folds.

Question6

```
tune_res=tune_grid(wf,resamples=folds,grid=grid)
```

```
autoplot(tune_res)
```



What do you notice? Do larger or smaller values of penalty and mixture produce better accuracy and ROC AUC?

Smaller values of penalty and mixture produce better accuracy and ROC AUC. Mixture values have no strong impact when penalty is large or small, but do have better accuracy and ROC AUC for mid-range penalty.

Question7

```
best=select_best(tune_res,metric="roc_auc")
final=finalize_workflow(wf,best)

final_fit=fit(final,data=train)

augment(final_fit,new_data=test) %>%
  accuracy(truth=type_1,estimate=.pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy multiclass 0.364
```


The performance on the testing set is pretty bad.

Question8


```
augment(final_fit,new_data=test) %>%  
  roc_auc(truth=type_1,estimate=.pred_Bug:.pred_Water)
```

```
## # A tibble: 1 x 3  
##   .metric .estimator .estimate  
##   <chr>   <chr>      <dbl>  
## 1 roc_auc hand_till      0.688
```

```
augment(final_fit,new_data=test) %>%  
  roc_curve(truth=type_1,estimate=.pred_Bug:.pred_Water) %>%  
  autoplot()
```



```
augment(final_fit,new_data=test) %>%  
  conf_mat(truth=type_1,estimate=.pred_class) %>%  
  autoplot(type="heatmap")
```

First of all, our model doesn't perform well. From the ROC curves, we can tell that normal type is the model best at predicting, which is also proved by the heatmap of the confusion matrix, followed by bug type and psychic type. Water type is the worst according to ROC curve, but heatmap shows different result. This might because water type has the most observations. The reason behind the poor performance of our model might be not enough observations. With less than 100 observations per type on average, it can be hard to achieve an accurate model.